

Data Management*

Chapter 5

Anna Ziff

R Workflow for Economists

Contents

Built-in Functions	2
Import and Export	2
R Saved Objects	4
Select Variables	5
Rename and Create Variables	6
Filter Observations	7
Organize	8
Merge	9
tidyverse Functions	10
Import and Export: readr	10
Practice Exercises 5.1	13
Transform: dplyr	13
Select Variables	13
Rename and Create Variables	15
Filter Observations	17
Organize	19
Merge	20
Practice Exercises 5.2	23
Reshape: tidyr	23
Pipes	25
Practice Exercises 5.3	28
Further Reading	28
References	28

*Please contact anna.ziff@duke.edu if there are errors.

In this chapter, we will go through the R functions needed for data management. The built-in R functions are useful tools and it is important to know their syntax. There are several packages that are widely used that are helpful to work with larger data, produce cleaner code, and be more efficient in data management. The suite of packages called **tidyverse** is especially common.

Here are all the libraries you should install for this chapter. Most of these are packages in **tidyverse**.

```
library(dplyr)
library(ggplot2)
library(magrittr)
library(readr)
library(readxl)
library(stringr)
library(tidyr)
```

We will practice importing data. Go to the [Dropbox folder](#) with the example data. Download the entire folder to a convenient file on your computer and save the file path for use in the below notes.

Built-in Functions

Import and Export

Importing text files, including those files with extensions `.txt` and `.csv` can be done with the function `read.table()`. This function reads a file and creates a data frame. The function `read.csv()` is a wrapper meaning it implements the same command but sets some defaults optimized for `.csv` files.

```
df1 <- read.csv("gapminder.csv")
str(df1)
```

```
## 'data.frame': 197 obs. of 4 variables:
## $ country: chr "Afghanistan" "Albania" "Algeria" "Andorra" ...
## $ gdp : int 574 4520 4780 42100 3750 13300 10600 3920 55100 47800 ...
## $ gini : num 36.8 29 27.6 40 42.6 40 41.8 31.9 32.3 30.6 ...
## $ region : chr "Asia & Pacific" "Europe" "Arab States" "Europe" ...
```

```
head(df1) # Display the first 5 rows
```

```
##           country  gdp gini           region
## 1      Afghanistan  574 36.8      Asia & Pacific
## 2         Albania 4520 29.0             Europe
## 3         Algeria 4780 27.6          Arab States
## 4         Andorra 42100 40.0             Europe
## 5          Angola  3750 42.6             Africa
## 6 Antigua and Barbuda 13300 40.0 South/Latin America
```

This command does the exact same thing.

```
df2 <- read.table("gapminder.csv", header = TRUE, sep = ",")
str(df2)
```

```
## 'data.frame': 197 obs. of 4 variables:
## $ country: chr "Afghanistan" "Albania" "Algeria" "Andorra" ...
## $ gdp : int 574 4520 4780 42100 3750 13300 10600 3920 55100 47800 ...
## $ gini : num 36.8 29 27.6 40 42.6 40 41.8 31.9 32.3 30.6 ...
## $ region : chr "Asia & Pacific" "Europe" "Arab States" "Europe" ...
```

Here is an example of reading a `.txt` file with `read.delim()`. Note that we need to specify the delimiter, in this case a space. You will need to inspect your file to determine the delimiter.

```
df3 <- read.delim("gapminder.txt", sep = " ")
str(df3)
```

```
## 'data.frame':    197 obs. of  4 variables:
## $ country: chr  "Afghanistan" "Albania" "Algeria" "Andorra" ...
## $ gdp : int  574 4520 4780 42100 3750 13300 10600 3920 55100 47800 ...
## $ gini : num  36.8 29 27.6 40 42.6 40 41.8 31.9 32.3 30.6 ...
## $ region : chr  "Asia & Pacific" "Europe" "Arab States" "Europe" ...
```

These three functions have many arguments available to adjust how the data files are read. The argument `stringsAsFactors` is automatically set to `FALSE`. If it is set to `TRUE`, then variables with character strings are read in as factors.

```
df4 <- read.csv("gapminder.csv", stringsAsFactors = TRUE)
str(df4)
```

```
## 'data.frame':    197 obs. of  4 variables:
## $ country: Factor w/ 195 levels "Afghanistan",...: 1 2 3 4 5 6 7 8 9 10 ...
## $ gdp : int  574 4520 4780 42100 3750 13300 10600 3920 55100 47800 ...
## $ gini : num  36.8 29 27.6 40 42.6 40 41.8 31.9 32.3 30.6 ...
## $ region : Factor w/ 7 levels "Africa","Arab States",...: 3 4 2 4 1 7 7 4 3 4 ...
```

You can specify the classes of all the columns using the argument `colClasses`. This is especially useful if the dataset is larger as it means that R does not need to determine the classes itself.

```
df5 <- read.csv("gapminder.csv",
               colClasses = c("character", "integer", "double", "factor"))
str(df5)
```

```
## 'data.frame':    197 obs. of  4 variables:
## $ country: chr  "Afghanistan" "Albania" "Algeria" "Andorra" ...
## $ gdp : int  574 4520 4780 42100 3750 13300 10600 3920 55100 47800 ...
## $ gini : num  36.8 29 27.6 40 42.6 40 41.8 31.9 32.3 30.6 ...
## $ region : Factor w/ 7 levels "Africa","Arab States",...: 3 4 2 4 1 7 7 4 3 4 ...
```

Column names (or variable names) and row names can be set while reading the file as well.

```
df6 <- read.csv("gapminder.csv",
               col.names = c("Country", "GDP", "GiniIndex", "Region"))
# row.names for rows
str(df6)
```

```
## 'data.frame':    197 obs. of  4 variables:
## $ Country : chr  "Afghanistan" "Albania" "Algeria" "Andorra" ...
## $ GDP : int  574 4520 4780 42100 3750 13300 10600 3920 55100 47800 ...
## $ GiniIndex: num  36.8 29 27.6 40 42.6 40 41.8 31.9 32.3 30.6 ...
## $ Region : chr  "Asia & Pacific" "Europe" "Arab States" "Europe" ...
```

If you just want to get a sense of what types of variables a dataset contains, you can use the `nrows` argument to read in very few rows. This is especially helpful with larger datasets.

```
checkcols <- read.csv("gapminder.csv",
                     nrows = 3)
checkcols
```

```
##      country  gdp gini      region
## 1 Afghanistan 574 36.8 Asia & Pacific
## 2   Albania 4520 29.0      Europe
## 3   Algeria 4780 27.6   Arab States
```

The built-in functions to export data are very similar to those to import data. Again, `write.table()` is the general function with `write.csv()` being a wrapper for different file types. Here is an example data frame that we will export.

```
df <- data.frame(id = seq(1:50),
                 v1 = rnorm(50, mean = 10, sd = 2),
                 v2 = rbinom(50, size = 1, prob = 0.5),
                 v3 = c(TRUE, FALSE),
                 v4 = c("Group 1", "Group 2", "Group 3", "Group 4", "Group 5"))
head(df)
```

```
##   id      v1 v2    v3     v4
## 1  1  9.718255 0 TRUE Group 1
## 2  2  9.689573 0 FALSE Group 2
## 3  3 10.072202 1 TRUE  Group 3
## 4  4 13.597885 1 FALSE Group 4
## 5  5 10.897250 1 TRUE  Group 5
## 6  6  9.775413 0 FALSE Group 1
```

Before exporting, make sure the correct directory is set. Remember you can use `getwd()` to check and `setwd()` to change the directory.

The function `write.csv()` exports a comma-delimited text file. You need to specify the object to be saved and the name of the file. The argument `row.names` determines whether the row names are exported as well. Unless you have custom row names, it is useful to set this argument to `FALSE`.

```
write.csv(df, file = "df_csv.csv", row.names = FALSE)
```

For greater generality, `write.table()` is available.

```
write.table(df, file = "df_table.txt", sep = "\t")
```

If you want to read Excel files, you will need an external package. A good option is the package `readxl` to access the function `read_excel()`. This package relies on tibbles just like `readr` (discussed below).

```
tib4 <- read_excel("gapminder.xlsx")
head(tib4)
```

```
## # A tibble: 6 x 4
##   country      gdp    gini region
##   <chr>      <chr> <dbl> <chr>
## 1 Afghanistan    574   36.8 Asia & Pacific
## 2 Albania        4520   29 Europe
## 3 Algeria        4780  27.6 Arab States
## 4 Andorra       42100  40 Europe
## 5 Angola        3750  42.6 Africa
## 6 Antigua and Barbuda 13300  40 South/Latin America
```

Other packages that allow you to read and write Excel files include `xlsx` and `r2excel`.

There are other packages that allow you to import and export datasets in other formats. For example, the `foreign` package allows for data files from SPSS, SAS, and STATA.

R Saved Objects

There are R-specific data formats to save the environment or components of it. To save the entire environment, use the `.RData` format.

```
ids <- 1:100
verbose_sqrt <- function(num) {
```

```

if (num >= 0) {
  return(sqrt(num))
} else {
  return("Negative number input.")
}
}
save(ids, verbose_sqrt, file = "workspace.RData")

```

This file includes both the objects and the names of the objects. You can directly load `.RData` and the workspace is populated. If you only want to save one object, you can use `.rds` files instead. These do not save the object's name. They are very memory-efficient (similar to saving a zipped file).

```
head(df)
```

```

##   id      v1 v2    v3      v4
## 1  1  9.718255 0 TRUE Group 1
## 2  2  9.689573 0 FALSE Group 2
## 3  3 10.072202 1 TRUE Group 3
## 4  4 13.597885 1 FALSE Group 4
## 5  5 10.897250 1 TRUE Group 5
## 6  6  9.775413 0 FALSE Group 1

```

```
saveRDS(df, "dataframe.rds")
```

Importing these objects is done as follows.

```

load("workspace.RData") # Imports objects and names
mydf <- readRDS("dataframe.rds") # Imports one object assigned to mydf

```

Select Variables

```

df <- read.csv("gapminder_large.csv")
str(df)

```

```

## 'data.frame':    195 obs. of  21 variables:
## $ country      : chr  "Afghanistan" "Albania" "Algeria" "Andorra" ...
## $ gdp_2015      : int   574 4520 4780 42100 3750 13300 10600 3920 55100 47800 ...
## $ gini_2015     : num   36.8 29 27.6 40 42.6 40 41.8 31.9 32.3 30.6 ...
## $ region       : chr   "Asia & Pacific" "Europe" "Arab States" "Europe" ...
## $ co2_2015      : num   0.262 1.6 3.8 5.97 1.22 5.84 4.64 1.65 16.8 7.7 ...
## $ co2_2016      : num   0.245 1.57 3.64 6.07 1.18 5.9 4.6 1.76 17 7.7 ...
## $ co2_2017      : num   0.247 1.61 3.56 6.27 1.14 5.89 4.55 1.7 17 7.94 ...
## $ co2_2018      : num   0.254 1.59 3.69 6.12 1.12 5.88 4.41 1.89 16.9 7.75 ...
## $ cpi_2012      : int    8 33 34 NA 22 NA 35 34 85 69 ...
## $ cpi_2013      : int    8 31 36 NA 23 NA 34 36 81 69 ...
## $ cpi_2014      : int   12 33 36 NA 19 NA 34 37 80 72 ...
## $ cpi_2015      : int   11 36 36 NA 15 NA 32 35 79 76 ...
## $ cpi_2016      : int   15 39 34 NA 18 NA 36 33 79 75 ...
## $ cpi_2017      : int   15 38 33 NA 19 NA 39 35 77 75 ...
## $ lifeexp_2012   : num   60.8 77.8 76.8 82.4 61.3 76.7 76 74.7 82.5 81 ...
## $ lifeexp_2013   : num   61.3 77.9 76.9 82.5 61.9 76.8 76.1 75.2 82.6 81.2 ...
## $ lifeexp_2014   : num   61.2 77.9 77 82.5 62.8 76.8 76.4 75.3 82.5 81.4 ...
## $ lifeexp_2015   : num   61.2 78 77.1 82.6 63.3 76.9 76.5 75.3 82.5 81.5 ...
## $ lifeexp_2016   : num   61.2 78.1 77.4 82.7 63.8 77 76.5 75.4 82.5 81.7 ...
## $ lifeexp_2017   : num   63.4 78.2 77.7 82.7 64.2 77 76.7 75.6 82.4 81.8 ...
## $ lifeexp_2018   : num   63.7 78.3 77.9 NA 64.6 77.2 76.8 75.8 82.5 81.9 ...

```

The built-in functions import data as data frames. Chapter 1 discusses how to select variables (columns). Here is a small review.

```
df[, 1:3]
df[, c(2, 4)]
df[, "cpi_2017"]
df[, c("lifeexp_2012", "cpi_2016")]
df[, c("country", "region")]
df[1:3]
df$gini_2015
```

Rename and Create Variables

The names of a data frame can be access with `names()`. This is an attribute of the data frame and can be used to rename all the variables this way.

```
names(df)

## [1] "country"      "gdp_2015"      "gini_2015"      "region"         "co2_2015"
## [6] "co2_2016"      "co2_2017"      "co2_2018"      "cpi_2012"       "cpi_2013"
## [11] "cpi_2014"      "cpi_2015"      "cpi_2016"      "cpi_2017"       "lifeexp_2012"
## [16] "lifeexp_2013" "lifeexp_2014" "lifeexp_2015" "lifeexp_2016"  "lifeexp_2017"
## [21] "lifeexp_2018"

names(df) <- paste0("var", 1:length(names(df)))
names(df)

## [1] "var1" "var2" "var3" "var4" "var5" "var6" "var7" "var8" "var9"
## [10] "var10" "var11" "var12" "var13" "var14" "var15" "var16" "var17" "var18"
## [19] "var19" "var20" "var21"
```

An alternative is to use the function `setNames()`. This function can also be used for other data structures besides data frames, such as vectors.

```
vnames <- c("country", "gdp_2015", "gini_2015", "region",
            "co2_2015", "co2_2016", "co2_2017", "co2_2018",
            "cpi_2012", "cpi_2013", "cpi_2014", "cpi_2015",
            "cpi_2016", "cpi_2017", "lifeexp_2012", "lifeexp_2013",
            "lifeexp_2014", "lifeexp_2015", "lifeexp_2016", "lifeexp_2017",
            "lifeexp_2018")
df <- setNames(df, vnames)
names(df)

## [1] "country"      "gdp_2015"      "gini_2015"      "region"         "co2_2015"
## [6] "co2_2016"      "co2_2017"      "co2_2018"      "cpi_2012"       "cpi_2013"
## [11] "cpi_2014"      "cpi_2015"      "cpi_2016"      "cpi_2017"       "lifeexp_2012"
## [16] "lifeexp_2013" "lifeexp_2014" "lifeexp_2015" "lifeexp_2016"  "lifeexp_2017"
## [21] "lifeexp_2018"
```

It is also possible to rename a subset of the variables.

```
names(df)[1] <- "COUNTRY"
names(df)

## [1] "COUNTRY"      "gdp_2015"      "gini_2015"      "region"         "co2_2015"
## [6] "co2_2016"      "co2_2017"      "co2_2018"      "cpi_2012"       "cpi_2013"
## [11] "cpi_2014"      "cpi_2015"      "cpi_2016"      "cpi_2017"       "lifeexp_2012"
## [16] "lifeexp_2013" "lifeexp_2014" "lifeexp_2015" "lifeexp_2016"  "lifeexp_2017"
## [21] "lifeexp_2018"
```

```
names(df)[2:3] <- c("GDP", "GINI")
names(df)
```

```
## [1] "COUNTRY"      "GDP"          "GINI"         "region"       "co2_2015"
## [6] "co2_2016"     "co2_2017"     "co2_2018"     "cpi_2012"     "cpi_2013"
## [11] "cpi_2014"     "cpi_2015"     "cpi_2016"     "cpi_2017"     "lifeexp_2012"
## [16] "lifeexp_2013" "lifeexp_2014" "lifeexp_2015" "lifeexp_2016" "lifeexp_2017"
## [21] "lifeexp_2018"
```

Creating new variables can be done with `cbind()` as discussed in chapter 1.

```
random1 <- rnorm(dim(df)[1])
head(random1)
```

```
## [1] -0.9034616 -0.9315785  0.7420078 -0.5748372  0.5144001 -0.5338265
```

```
df <- cbind(df, random1)
df[1:5, c("COUNTRY", "random1")]
```

```
##      COUNTRY  random1
## 1 Afghanistan -0.9034616
## 2   Albania -0.9315785
## 3   Algeria  0.7420078
## 4   Andorra -0.5748372
## 5    Angola  0.5144001
```

This method has the advantage that it can be used to add more than one variable at a time.

```
random2 <- runif(dim(df)[1])
random3 <- rexp(dim(df)[1])
df <- cbind(df, random2, random3)
df[1:5, c("COUNTRY", "random2", "random3")]
```

```
##      COUNTRY  random2  random3
## 1 Afghanistan 0.2628448 0.7706266
## 2   Albania 0.8769630 0.8532965
## 3   Algeria 0.9611621 0.2807853
## 4   Andorra 0.6933212 1.9286870
## 5    Angola 0.9711144 0.2093739
```

The following shortcut is helpful to create one variable at a time.

```
df$random4 <- df$random3^2
df[1:5, c("COUNTRY", "random4")]
```

```
##      COUNTRY  random4
## 1 Afghanistan 0.59386536
## 2   Albania 0.72811496
## 3   Algeria 0.07884039
## 4   Andorra 3.71983369
## 5    Angola 0.04383745
```

Filter Observations

Filtering observations can be done by row name or number, as shown in chapter 1.

```
df[1:3, ]
df[c(3, 40), ]
df[c("4", "17"), ]
```

```
df[!c(1:190), ]
df[-c(1:190), ]
```

Filtering can also be done using logical statements.

```
df[df$random2 >= 1, ]
```

```
## [1] COUNTRY      GDP      GINI      region    co2_2015
## [6] co2_2016    co2_2017    co2_2018    cpi_2012    cpi_2013
## [11] cpi_2014    cpi_2015    cpi_2016    cpi_2017    lifeexp_2012
## [16] lifeexp_2013 lifeexp_2014 lifeexp_2015 lifeexp_2016 lifeexp_2017
## [21] lifeexp_2018 random1      random2      random3      random4
## <0 rows> (or 0-length row.names)
```

```
df[df$random2 >= 1 & df$random3 <= 0.5, ]
```

```
## [1] COUNTRY      GDP      GINI      region    co2_2015
## [6] co2_2016    co2_2017    co2_2018    cpi_2012    cpi_2013
## [11] cpi_2014    cpi_2015    cpi_2016    cpi_2017    lifeexp_2012
## [16] lifeexp_2013 lifeexp_2014 lifeexp_2015 lifeexp_2016 lifeexp_2017
## [21] lifeexp_2018 random1      random2      random3      random4
## <0 rows> (or 0-length row.names)
```

```
subset(df, df$random3 <= 0.05)[, c("COUNTRY", "random3")]
```

```
##          COUNTRY      random3
## 17          Belgium 0.017757062
## 35           Chile 0.007570992
## 46 Czech Republic 0.021670612
## 64          Georgia 0.030343974
## 65          Germany 0.043363370
## 98 Liechtenstein 0.014375120
## 99          Lithuania 0.036783109
## 100 Luxembourg 0.040162183
## 119          Namibia 0.006381497
## 126          Nigeria 0.011836921
## 135          Paraguay 0.018873897
## 144           Samoa 0.028333847
## 147 Saudi Arabia 0.023012085
## 178           Tuvalu 0.029306085
```

The `which()` function returns the row numbers that are being filtered.

```
which(df$random3 <= 0.05)
```

```
## [1] 17 35 46 64 65 98 99 100 119 126 135 144 147 178
```

Organize

Sorting can be done by one or more columns. Note that even though the rows are re-ordered, the original row names remain.

```
dforder1 <- order(df$GINI)
head(df[dforder1, c("COUNTRY", "GINI")])
```

```
##          COUNTRY GINI
## 180          Ukraine 24.8
## 154          Slovenia 25.6
```



```
## 46    Czech Republic 26.0
## 153 Slovak Republic 26.7
## 16      Belarus 26.9
## 87      Kazakhstan 26.9
```

```
dforder2 <- order(df$region, df$GINI)
head(df[dforder2, c("COUNTRY", "region", "GINI")])
```

```
##              COUNTRY region GINI
## 146 Sao Tome and Principe Africa 30.8
## 105              Mali Africa 33.0
## 96      Liberia Africa 33.3
## 70      Guinea Africa 33.7
## 151      Sierra Leone Africa 34.0
## 125              Niger Africa 34.1
```

Merge

As discussed in chapter 1, `rbind()` can be used to append additional observations. If using this approach, it is better to transform the new row(s) into a data frame. This will help avoid silently changing a variable type.

```
df1 <- df[1:98, ]
df2 <- df[99:195, ]
rbind(df1, df2)
```

An even more robust approach is to use the `merge()` function. This allows for the two data frames to have different variables and similar observations. As long as there is at least one variable common to both data frames, they can be merged. Here is a very simple example.

```
df1 <- df[1:5, c("COUNTRY", "region")]
df2 <- df[1:7, c("COUNTRY", "GDP", "GINI")]
merge(df1, df2, by = "COUNTRY")
```

```
##      COUNTRY      region  GDP GINI
## 1 Afghanistan Asia & Pacific  574 36.8
## 2    Albania      Europe 4520 29.0
## 3    Algeria Arab States  4780 27.6
## 4    Andorra      Europe 42100 40.0
## 5     Angola      Africa  3750 42.6
```

Note that `df2` has 7 observations while `df1` only has 5. Yet, the output of the merge has 5 observations. This is because the arguments `all.x` and `all.y` are set to `FALSE` by default. This means that only rows that appear in both are present in the output. If we set `all.y = TRUE`, all the rows of `df2` are added with missing values for `region`.

```
merge(df1, df2, by = "COUNTRY", all.y = TRUE)
```

```
##      COUNTRY      region  GDP GINI
## 1 Afghanistan Asia & Pacific  574 36.8
## 2    Albania      Europe 4520 29.0
## 3    Algeria Arab States  4780 27.6
## 4    Andorra      Europe 42100 40.0
## 5     Angola      Africa  3750 42.6
## 6  Antigua and Barbuda    <NA> 13300 40.0
## 7    Argentina    <NA> 10600 41.8
```

If you want to keep all the rows in both data frames, the argument `all = TRUE` sets both `all.x = TRUE` and `all.y = TRUE`.

```
merge(df1, df2, by = "COUNTRY", all = TRUE)
```

```
##           COUNTRY      region  GDP GINI
## 1  Afghanistan Asia & Pacific   574 36.8
## 2      Albania      Europe  4520 29.0
## 3      Algeria  Arab States  4780 27.6
## 4      Andorra      Europe 42100 40.0
## 5       Angola      Africa  3750 42.6
## 6 Antigua and Barbuda      <NA> 13300 40.0
## 7      Argentina      <NA> 10600 41.8
```

Suppose the variable you are merging on has different names in the two data frames. The arguments `by.x` and `by.y` allow for you to specify both variables.

```
names(df1)[1] <- "country"
merge(df1, df2, by.x = "country", by.y = "COUNTRY")
```

```
##      country      region  GDP GINI
## 1 Afghanistan Asia & Pacific   574 36.8
## 2      Albania      Europe  4520 29.0
## 3      Algeria  Arab States  4780 27.6
## 4      Andorra      Europe 42100 40.0
## 5       Angola      Africa  3750 42.6
```

If the two data frames have different variables with the same name, the merge will not combine these columns. This even applies if the columns are different types.

```
df1 <- df[c("COUNTRY", "region", "GDP")]
df1$GDP <- as.character(df1$GDP) # GDP is now character in df1
merge(df1, df2, by = "COUNTRY")
```

```
##           COUNTRY      region GDP.x GDP.y GINI
## 1  Afghanistan Asia & Pacific   574   574 36.8
## 2      Albania      Europe  4520  4520 29.0
## 3      Algeria  Arab States  4780  4780 27.6
## 4      Andorra      Europe 42100 42100 40.0
## 5       Angola      Africa  3750  3750 42.6
## 6 Antigua and Barbuda South/Latin America 13300 13300 40.0
## 7      Argentina South/Latin America 10600 10600 41.8
```

tidyverse Functions

Hadley Wickham developed the idea behind a suite of packages that streamline data work called **tidyverse**. There are many packages in this suite that relate to different types of datasets and parts of the data process. This chapter goes through `dplyr`, `tidyr`, and `readr`.

Import and Export: readr

The functions in the `readr` package to read and write data are faster than the built-in functions. Apart from efficiency, they have another advantage in that they help ensure consistency in the imported data. For example, if there are spaces in the variable name, `read.csv()`, the built-in function, will automatically remove these. The `readr` function `read_csv()` will not remove them.

```
tib1 <- read_csv("gapminder.csv")
```

```
## Rows: 197 Columns: 4
```

```
## -- Column specification -----
## Delimiter: ","
## chr (2): country, region
## dbl (2): gdp, gini
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
str(tib1)

## spec_tbl_df [197 x 4] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ country: chr [1:197] "Afghanistan" "Albania" "Algeria" "Andorra" ...
## $ gdp : num [1:197] 574 4520 4780 42100 3750 13300 10600 3920 55100 47800 ...
## $ gini : num [1:197] 36.8 29 27.6 40 42.6 40 41.8 31.9 32.3 30.6 ...
## $ region : chr [1:197] "Asia & Pacific" "Europe" "Arab States" "Europe" ...
## - attr(*, "spec")=
## .. cols(
## .. country = col_character(),
## .. gdp = col_double(),
## .. gini = col_double(),
## .. region = col_character()
## .. )
## - attr(*, "problems")=<externalptr>
```

Immediately, you can see that the data structure is different. The package `readr`, and all the packages in the `tidyverse` suite, rely on a data structure called tibbles instead of data frames. The two main differences between tibbles and data frames are the following. More information on the differences is available [here](#).

- Unlike data frames, tibbles only show the first 10 rows and enough columns to fit on the screen. Each column is printed with its type.
- When subsetting, `[]` always returns another tibble and `[[]]` always returns a vector.

Just like in `read.csv()`, you can specify the columns.

```
tib2 <- read_csv("gapminder.csv",
  col_types = list(col_character(),
    col_integer(),
    col_double(),
    col_factor()))

str(tib2)

## spec_tbl_df [197 x 4] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ country: chr [1:197] "Afghanistan" "Albania" "Algeria" "Andorra" ...
## $ gdp : int [1:197] 574 4520 4780 42100 3750 13300 10600 3920 55100 47800 ...
## $ gini : num [1:197] 36.8 29 27.6 40 42.6 40 41.8 31.9 32.3 30.6 ...
## $ region : Factor w/ 7 levels "Asia & Pacific",...: 1 2 3 2 4 5 5 2 1 2 ...
## - attr(*, "spec")=
## .. cols(
## .. country = col_character(),
## .. gdp = col_integer(),
## .. gini = col_double(),
## .. region = col_factor(levels = NULL, ordered = FALSE, include_na = FALSE)
## .. )
## - attr(*, "problems")=<externalptr>
```

If you want to completely rename the columns, you can do so with the option `col_names`. You will just need to tell R to skip reading in the first line of the file.

```
tib3 <- read_csv("gapminder.csv", skip = 1,
                 col_names = c("V1", "V2", "V3", "V4"))

## Rows: 197 Columns: 4
## -- Column specification -----
## Delimiter: ","
## chr (2): V1, V4
## dbl (2): V2, V3
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
str(tib3)
```

```
## spec_tbl_df [197 x 4] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ V1: chr [1:197] "Afghanistan" "Albania" "Algeria" "Andorra" ...
## $ V2: num [1:197] 574 4520 4780 42100 3750 13300 10600 3920 55100 47800 ...
## $ V3: num [1:197] 36.8 29 27.6 40 42.6 40 41.8 31.9 32.3 30.6 ...
## $ V4: chr [1:197] "Asia & Pacific" "Europe" "Arab States" "Europe" ...
## - attr(*, "spec")=
## .. cols(
## .. V1 = col_character(),
## .. V2 = col_double(),
## .. V3 = col_double(),
## .. V4 = col_character()
## .. )
## - attr(*, "problems")=<externalptr>
```

The argument `n_max` determines the maximum number of lines that are read.

```
read_csv("gapminder.csv", n_max = 3)

## Rows: 3 Columns: 4
## -- Column specification -----
## Delimiter: ","
## chr (2): country, region
## dbl (2): gdp, gini
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

## # A tibble: 3 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Afghanistan  574   36.8 Asia & Pacific
## 2 Albania     4520   29   Europe
## 3 Algeria     4780  27.6 Arab States
```

The `readr` analogues to `read.table()` and `read.delim()` are `read_table()` and `read_delim()`. They have similar arguments as `read_csv()`. Reading in data files usually presents unexpected difficulties and complications, and the myriad of arguments available can help address any formatting issues automatically.

The write functions in `readr` are faster than the built-in functions and automatically omit row names.

```
write_csv(df, file = "df_csv_readr.csv")
```

Practice Exercises 5.1

1. Another way to list the column types is string shortcuts. For example "d" for double, "c" for character, etc. Check the documentation for `read_csv()`, and call in "gapminder.csv" with a character column, an integer column, a double column, and a factor column.
2. You can also easily skip columns with this shorthand. Why do you think this be useful? Call in "gapminder.csv" again skipping the Region column.

Transform: dplyr

The package `dplyr` includes functions that transform tibbles and data frames.

```
df <- read_csv("gapminder.csv")

## Rows: 197 Columns: 4
## -- Column specification -----
## Delimiter: ","
## chr (2): country, region
## dbl (2): gdp, gini
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
head(df)

## # A tibble: 6 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Afghanistan    574  36.8 Asia & Pacific
## 2 Albania        4520  29 Europe
## 3 Algeria        4780  27.6 Arab States
## 4 Andorra       42100  40 Europe
## 5 Angola         3750  42.6 Africa
## 6 Antigua and Barbuda 13300  40 South/Latin America
```

Select Variables

The general form of functions in `dplyr` involves identifying the data frame first and then specifying the options. To demonstrate, the function `select` chooses which variables.

```
select(tib1, country)

## # A tibble: 197 x 1
##   country
##   <chr>
## 1 Afghanistan
## 2 Albania
## 3 Algeria
## 4 Andorra
## 5 Angola
## 6 Antigua and Barbuda
## 7 Argentina
## 8 Armenia
## 9 Australia
## 10 Austria
## # ... with 187 more rows
```

Note that the original data frame is not changed. You will have to assign an object if you want to save this selection in an object.

```
head(tib1)
```

```
## # A tibble: 6 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Afghanistan    574  36.8 Asia & Pacific
## 2 Albania       4520  29   Europe
## 3 Algeria       4780  27.6 Arab States
## 4 Andorra      42100  40   Europe
## 5 Angola        3750  42.6 Africa
## 6 Antigua and Barbuda 13300  40   South/Latin America
```

There are several ways to select more than one variable. The last method used a helper, `starts_with()`. See the documentation for `select` for other helpers.

```
select(tib1, country, gdp)
select(tib1, gdp:gini)
select(tib1, -gdp)
select(tib1, -c(country, gini))
select(tib1, starts_with("g"))
```

Sometimes it is desirable to rename variables when selecting them. This is very convenient in `select`!

```
select(tib1, country_name = country)
```

```
## # A tibble: 197 x 1
##   country_name
##   <chr>
## 1 Afghanistan
## 2 Albania
## 3 Algeria
## 4 Andorra
## 5 Angola
## 6 Antigua and Barbuda
## 7 Argentina
## 8 Armenia
## 9 Australia
## 10 Austria
## # ... with 187 more rows
```

```
select(tib1, var = starts_with("g"))
```

```
## # A tibble: 197 x 2
##   var1 var2
##   <dbl> <dbl>
## 1    574  36.8
## 2   4520  29
## 3   4780  27.6
## 4  42100  40
## 5   3750  42.6
## 6  13300  40
## 7  10600  41.8
## 8   3920  31.9
## 9  55100  32.3
```

```
## 10 47800 30.6
## # ... with 187 more rows
```

Rename and Create Variables

If you want to rename variables without dropping any, use the function `rename`.

```
rename(tib1, country_name = country)
```

```
## # A tibble: 197 x 4
##   country_name      gdp  gini region
##   <chr>          <dbl> <dbl> <chr>
## 1 Afghanistan      574  36.8 Asia & Pacific
## 2 Albania           4520  29 Europe
## 3 Algeria           4780  27.6 Arab States
## 4 Andorra          42100  40 Europe
## 5 Angola            3750  42.6 Africa
## 6 Antigua and Barbuda 13300  40 South/Latin America
## 7 Argentina         10600  41.8 South/Latin America
## 8 Armenia           3920  31.9 Europe
## 9 Australia         55100  32.3 Asia & Pacific
## 10 Austria          47800  30.6 Europe
## # ... with 187 more rows
```

```
rename(tib1, country_name = country, gdp_percapita = gdp)
```

```
## # A tibble: 197 x 4
##   country_name      gdp_percapita  gini region
##   <chr>          <dbl> <dbl> <chr>
## 1 Afghanistan      574  36.8 Asia & Pacific
## 2 Albania           4520  29 Europe
## 3 Algeria           4780  27.6 Arab States
## 4 Andorra          42100  40 Europe
## 5 Angola            3750  42.6 Africa
## 6 Antigua and Barbuda 13300  40 South/Latin America
## 7 Argentina         10600  41.8 South/Latin America
## 8 Armenia           3920  31.9 Europe
## 9 Australia         55100  32.3 Asia & Pacific
## 10 Austria          47800  30.6 Europe
## # ... with 187 more rows
```

The function `mutate()` allows for new variables to be added to the data frame or existing variables to be modified without changing the other variables.

```
mutate(tib1, gdp_sq = gdp^2)
```

```
## # A tibble: 197 x 5
##   country      gdp  gini region      gdp_sq
##   <chr>          <dbl> <dbl> <chr>          <dbl>
## 1 Afghanistan      574  36.8 Asia & Pacific    329476
## 2 Albania           4520  29 Europe       20430400
## 3 Algeria           4780  27.6 Arab States   22848400
## 4 Andorra          42100  40 Europe      1772410000
## 5 Angola            3750  42.6 Africa       14062500
## 6 Antigua and Barbuda 13300  40 South/Latin America 176890000
## 7 Argentina         10600  41.8 South/Latin America 112360000
## 8 Armenia           3920  31.9 Europe       15366400
```

```
## 9 Australia          55100 32.3 Asia & Pacific    3036010000
## 10 Austria           47800 30.6 Europe          2284840000
## # ... with 187 more rows
```

```
mutate(tib1, row_id = 1:length(country))
```

```
## # A tibble: 197 x 5
##   country      gdp  gini region    row_id
##   <chr>      <dbl> <dbl> <chr>    <int>
## 1 Afghanistan    574  36.8 Asia & Pacific      1
## 2 Albania        4520  29 Europe            2
## 3 Algeria        4780  27.6 Arab States      3
## 4 Andorra       42100  40 Europe            4
## 5 Angola         3750  42.6 Africa           5
## 6 Antigua and Barbuda 13300  40 South/Latin America  6
## 7 Argentina     10600  41.8 South/Latin America  7
## 8 Armenia        3920  31.9 Europe           8
## 9 Australia     55100  32.3 Asia & Pacific     9
## 10 Austria      47800  30.6 Europe          10
## # ... with 187 more rows
```

```
mutate(tib1, gdp_large = ifelse(gdp >= 25000, TRUE, FALSE))
```

```
## # A tibble: 197 x 5
##   country      gdp  gini region    gdp_large
##   <chr>      <dbl> <dbl> <chr>    <lgl>
## 1 Afghanistan    574  36.8 Asia & Pacific FALSE
## 2 Albania        4520  29 Europe    FALSE
## 3 Algeria        4780  27.6 Arab States FALSE
## 4 Andorra       42100  40 Europe    TRUE
## 5 Angola         3750  42.6 Africa  FALSE
## 6 Antigua and Barbuda 13300  40 South/Latin America FALSE
## 7 Argentina     10600  41.8 South/Latin America FALSE
## 8 Armenia        3920  31.9 Europe  FALSE
## 9 Australia     55100  32.3 Asia & Pacific TRUE
## 10 Austria      47800  30.6 Europe  TRUE
## # ... with 187 more rows
```

If you want to create a new variable and drop the other variables, use the function `transmute()`.

```
transmute(tib1, gini_small = ifelse(gini <= 40, TRUE, FALSE))
```

```
## # A tibble: 197 x 1
##   gini_small
##   <lgl>
## 1 TRUE
## 2 TRUE
## 3 TRUE
## 4 TRUE
## 5 FALSE
## 6 TRUE
## 7 FALSE
## 8 TRUE
## 9 TRUE
## 10 TRUE
## # ... with 187 more rows
```


Filter Observations

The function `select` allows you to choose which variables (columns) are included in your data. The function `filter` allows you choose which observations (rows) are included in your data.

```
filter(tib1, region == "North America")
```

```
## # A tibble: 2 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Canada      50300  31.7 North America
## 2 United States 52100  41.3 North America
```

```
filter(tib1, is.na(gdp))
```

```
## # A tibble: 8 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Djibouti      NA  44.1 Arab States
## 2 Eritrea       NA   40  Africa
## 3 Liechtenstein NA   40  Europe
## 4 Venezuela     NA  46.9 South/Latin America
## 5 Holy See      NA   40  Europe
## 6 North Korea   NA   37  Asia & Pacific
## 7 Somalia       NA   48  Arab States
## 8 Syria         NA  35.2 Asia & Pacific
```

```
filter(tib1, gdp > 25000 & region != "Europe")
```

```
## # A tibble: 12 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Australia    55100  32.3 Asia & Pacific
## 2 Bahamas      27500  43.7 South/Latin America
## 3 Brunei       32900   40  Asia & Pacific
## 4 Canada       50300  31.7 North America
## 5 Japan        47100  32.1 Asia & Pacific
## 6 Kuwait       36000   40  Middle east
## 7 New Zealand  36800  34.5 Asia & Pacific
## 8 Qatar        65100   40  Middle east
## 9 Singapore    54000  40.9 Asia & Pacific
## 10 South Korea  26100  31.6 Asia & Pacific
## 11 United Arab Emirates 40200   40  Middle east
## 12 United States 52100  41.3 North America
```

```
filter(tib1, region %in% c("North America", "Middle east"))
```

```
## # A tibble: 14 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Canada      50300  31.7 North America
## 2 Egypt       2700  31.2 Middle east
## 3 Iran        6070  38.5 Middle east
## 4 Iraq        5300  29.5 Middle east
## 5 Jordan      3310  33.7 Middle east
## 6 Kuwait     36000   40  Middle east
## 7 Lebanon     6490  31.8 Middle east
```

```
## 8 Libya          5900 40 Middle east
## 9 Oman           16200 40 Middle east
## 10 Qatar          65100 40 Middle east
## 11 Saudi Arabia   21400 40 Middle east
## 12 United Arab Emirates 40200 40 Middle east
## 13 United States   52100 41.3 North America
## 14 Yemen           785 36.7 Middle east
```

To select rows based on the number index, use `slice`.

```
slice(tib1, 32:37)
```

```
## # A tibble: 6 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Cape Verde    3410  47.2 Africa
## 2 Central African Republic 347  56.2 Africa
## 3 Chad          957  43.3 Africa
## 4 Chile        14700  47.5 South/Latin America
## 5 China         6500  39.4 Asia & Pacific
## 6 Colombia      7580  51.7 South/Latin America
```

The function `distinct()` filters out duplicated rows.

```
distinct(tib1)
```

```
## # A tibble: 195 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Afghanistan    574  36.8 Asia & Pacific
## 2 Albania        4520  29 Europe
## 3 Algeria        4780  27.6 Arab States
## 4 Andorra       42100  40 Europe
## 5 Angola         3750  42.6 Africa
## 6 Antigua and Barbuda 13300  40 South/Latin America
## 7 Argentina     10600  41.8 South/Latin America
## 8 Armenia        3920  31.9 Europe
## 9 Australia     55100  32.3 Asia & Pacific
## 10 Austria       47800  30.6 Europe
## # ... with 185 more rows
```

```
filter(tib1, duplicated(tib1)) # Check which observations are duplicated
```

```
## # A tibble: 2 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Norway    90000  27.1 Europe
## 2 Suriname  8460  61 South/Latin America
```

The function `slice_sample()` randomly selects rows.

```
slice_sample(tib1, n = 4)
```

```
## # A tibble: 4 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Guinea-Bissau    574  50.7 Africa
## 2 Dominica        6890  40 South/Latin America
```

```
## 3 Portugal      22000  35.7 Europe
## 4 Uganda        900   41.9 Africa
```

```
slice_sample(tib1, prop = 0.03)
```

```
## # A tibble: 5 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Suriname    8460   61  South/Latin America
## 2 Ethiopia    483   37.7 Africa
## 3 United Kingdom 42000  33.4 Europe
## 4 South Sudan  731   45   Africa
## 5 Tajikistan  936   33.7 Asia & Pacific
```

Organize

The functions so far produce data frames that explicitly differ from the inputted data frame. There are some silent functions that change the underlying structure without changing the outputted data frame. The function `group_by()` is an example of these silent functions. It groups the data based on the values of a set of variables. It makes most sense to group by categorical variables. The only difference is that now it says Groups: region [7].

```
group_tib1 <- group_by(tib1, region)
group_tib1
```

```
## # A tibble: 197 x 4
## # Groups:   region [7]
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Afghanistan    574  36.8 Asia & Pacific
## 2 Albania        4520  29   Europe
## 3 Algeria        4780  27.6 Arab States
## 4 Andorra       42100  40   Europe
## 5 Angola         3750  42.6 Africa
## 6 Antigua and Barbuda 13300  40   South/Latin America
## 7 Argentina     10600  41.8 South/Latin America
## 8 Armenia        3920  31.9 Europe
## 9 Australia     55100  32.3 Asia & Pacific
## 10 Austria       47800  30.6 Europe
## # ... with 187 more rows
```

Ungrouping the data is another silent function and it removes this underlying grouping.

```
ungroup(group_tib1)
```

```
## # A tibble: 197 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Afghanistan    574  36.8 Asia & Pacific
## 2 Albania        4520  29   Europe
## 3 Algeria        4780  27.6 Arab States
## 4 Andorra       42100  40   Europe
## 5 Angola         3750  42.6 Africa
## 6 Antigua and Barbuda 13300  40   South/Latin America
## 7 Argentina     10600  41.8 South/Latin America
## 8 Armenia        3920  31.9 Europe
## 9 Australia     55100  32.3 Asia & Pacific
```

```
## 10 Austria          47800  30.6 Europe
## # ... with 187 more rows
```

The function `arrange` sorts the data based on the rank order of a set of variables. Adding `desc()` changes the rank-order to descending.

```
arrange(tib1, gini)
```

```
## # A tibble: 197 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Ukraine      2830  24.8 Europe
## 2 Slovenia    23800  25.6 Europe
## 3 Czech Republic 21400  26   Europe
## 4 Slovak Republic 18900  26.7 Europe
## 5 Belarus      6380  26.9 Europe
## 6 Kazakhstan   10600  26.9 Asia & Pacific
## 7 Moldova      2950  27   Europe
## 8 Finland     45600  27.1 Europe
## 9 Norway      90000  27.1 Europe
## 10 Norway      90000  27.1 Europe
## # ... with 187 more rows
```

```
arrange(tib1, desc(region), gini)
```

```
## # A tibble: 197 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Antigua and Barbuda 13300  40   South/Latin America
## 2 Dominica      6890  40   South/Latin America
## 3 Grenada       8190  40   South/Latin America
## 4 St. Kitts and Nevis 16700  40   South/Latin America
## 5 St. Vincent and the Grenadines 6580  40   South/Latin America
## 6 Uruguay      13900 40.1 South/Latin America
## 7 El Salvador     3310 41.1 South/Latin America
## 8 Trinidad and Tobago 16800 41.3 South/Latin America
## 9 Argentina     10600 41.8 South/Latin America
## 10 St. Lucia      8490 42.6 South/Latin America
## # ... with 187 more rows
```

Merge

Merging data frames is useful when there are several data frames with similar observations but different variables. To demonstrate the join functions in `dplyr`, we have two datasets. One is the population of all countries and the other is the population of all countries that begin with “A.” Neither of these datasets have duplicates.

```
pop <- read_csv("population.csv")
```

```
## Rows: 195 Columns: 2
## -- Column specification -----
## Delimiter: ","
## chr (1): country
## dbl (1): population
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
popA <- read_csv("population_A.csv")
```

```
## Rows: 11 Columns: 2
## -- Column specification -----
## Delimiter: ","
## chr (1): country
## dbl (1): population
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

The different join functions relate to which observations are kept. In `full_join()`, all observations in the two data frames are kept, even if there are unmatched observations. The argument `by` indicates which variable on which to match.

```
full_join(tib1, pop, by = "country")
```

```
## # A tibble: 197 x 5
##   country      gdp  gini region      population
##   <chr>      <dbl> <dbl> <chr>      <dbl>
## 1 Afghanistan    574  36.8 Asia & Pacific 34400000
## 2 Albania        4520  29 Europe      2890000
## 3 Algeria        4780  27.6 Arab States 39700000
## 4 Andorra       42100  40 Europe       78000
## 5 Angola         3750  42.6 Africa    27900000
## 6 Antigua and Barbuda 13300  40 South/Latin America 93600
## 7 Argentina      10600  41.8 South/Latin America 43100000
## 8 Armenia        3920  31.9 Europe     2930000
## 9 Australia      55100  32.3 Asia & Pacific 23900000
## 10 Austria       47800  30.6 Europe     8680000
## # ... with 187 more rows
```

The function `inner_join()` only keeps observations that are present in both data frames. In this case, that is only countries that begin with “A.”

```
inner_join(tib1, popA, by = "country")
```

```
## # A tibble: 11 x 5
##   country      gdp  gini region      population
##   <chr>      <dbl> <dbl> <chr>      <dbl>
## 1 Afghanistan    574  36.8 Asia & Pacific 34400000
## 2 Albania        4520  29 Europe      2890000
## 3 Algeria        4780  27.6 Arab States 39700000
## 4 Andorra       42100  40 Europe       78000
## 5 Angola         3750  42.6 Africa    27900000
## 6 Antigua and Barbuda 13300  40 South/Latin America 93600
## 7 Argentina      10600  41.8 South/Latin America 43100000
## 8 Armenia        3920  31.9 Europe     2930000
## 9 Australia      55100  32.3 Asia & Pacific 23900000
## 10 Austria       47800  30.6 Europe     8680000
## 11 Azerbaijan     6060  32.4 Asia & Pacific 9620000
```

The function `left_join()` only keeps from the data frame in the left argument (`tib1` in this case).

```
left_join(tib1, popA, by = "country")
```

```
## # A tibble: 197 x 5
```

```
##   country      gdp  gini region      population
##   <chr>      <dbl> <dbl> <chr>      <dbl>
## 1 Afghanistan    574  36.8 Asia & Pacific  34400000
## 2 Albania        4520  29   Europe      2890000
## 3 Algeria        4780  27.6 Arab States  39700000
## 4 Andorra       42100  40   Europe       78000
## 5 Angola         3750  42.6 Africa     27900000
## 6 Antigua and Barbuda 13300  40   South/Latin America  93600
## 7 Argentina      10600  41.8 South/Latin America  43100000
## 8 Armenia        3920  31.9 Europe     2930000
## 9 Australia      55100  32.3 Asia & Pacific  23900000
## 10 Austria       47800  30.6 Europe     8680000
## # ... with 187 more rows
```

The function `right_join()` is the same except it only keeps the observations from the data frame in the right argument.

```
right_join(tib1, popA, by = "country")
```

```
## # A tibble: 11 x 5
##   country      gdp  gini region      population
##   <chr>      <dbl> <dbl> <chr>      <dbl>
## 1 Afghanistan    574  36.8 Asia & Pacific  34400000
## 2 Albania        4520  29   Europe      2890000
## 3 Algeria        4780  27.6 Arab States  39700000
## 4 Andorra       42100  40   Europe       78000
## 5 Angola         3750  42.6 Africa     27900000
## 6 Antigua and Barbuda 13300  40   South/Latin America  93600
## 7 Argentina      10600  41.8 South/Latin America  43100000
## 8 Armenia        3920  31.9 Europe     2930000
## 9 Australia      55100  32.3 Asia & Pacific  23900000
## 10 Austria       47800  30.6 Europe     8680000
## 11 Azerbaijan     6060  32.4 Asia & Pacific  9620000
```

The function `semi_join()` keeps all rows in `tib1` that have a match in `popA`.

```
semi_join(tib1, popA, by = "country")
```

```
## # A tibble: 11 x 4
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Afghanistan    574  36.8 Asia & Pacific
## 2 Albania        4520  29   Europe
## 3 Algeria        4780  27.6 Arab States
## 4 Andorra       42100  40   Europe
## 5 Angola         3750  42.6 Africa
## 6 Antigua and Barbuda 13300  40   South/Latin America
## 7 Argentina      10600  41.8 South/Latin America
## 8 Armenia        3920  31.9 Europe
## 9 Australia      55100  32.3 Asia & Pacific
## 10 Austria       47800  30.6 Europe
## 11 Azerbaijan     6060  32.4 Asia & Pacific
```

The function `anti_join()` keeps all rows in `tib1` that do not have a match in `popA`.

```
anti_join(tib1, popA, by = "country")
```

```
## # A tibble: 186 x 4
```

```
##   country      gdp  gini region
##   <chr>      <dbl> <dbl> <chr>
## 1 Bahamas    27500  43.7 South/Latin America
## 2 Bahrain     22400   40   Arab States
## 3 Bangladesh  1000   32.3 Asia & Pacific
## 4 Barbados    15800  43.8 South/Latin America
## 5 Belarus      6380  26.9 Europe
## 6 Belgium     45500  27.8 Europe
## 7 Belize       4300  53.3 South/Latin America
## 8 Benin        1130  46.9 Africa
## 9 Bhutan       2780   38   Asia & Pacific
## 10 Bolivia     2360  46.3 South/Latin America
## # ... with 176 more rows
```

Practice Exercises 5.2

1. Before running this code, what do you think the output will be? Check to see if you were right!

```
anti_join(popA, tib1, by = "country")
```

Reshape: tidyr

The tidyr package provides an efficient way to reshape and reformat data.

```
tib1 <- read_csv("gapminder_large.csv")
```

```
## Rows: 195 Columns: 21
## -- Column specification -----
## Delimiter: ","
## chr (2): country, region
## dbl (19): gdp_2015, gini_2015, co2_2015, co2_2016, co2_2017, co2_2018, cpi_2...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
head(tib1)
```

```
## # A tibble: 6 x 21
##   country gdp_2015 gini_2015 region co2_2015 co2_2016 co2_2017 co2_2018 cpi_2012
##   <chr>      <dbl>    <dbl> <chr>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 Afghan~      574      36.8 Asia ~    0.262    0.245    0.247    0.254      8
## 2 Albania    4520      29   Europe    1.6      1.57    1.61    1.59     33
## 3 Algeria    4780      27.6 Arab ~    3.8      3.64    3.56    3.69     34
## 4 Andorra   42100      40   Europe    5.97     6.07    6.27    6.12     NA
## 5 Angola     3750      42.6 Africa    1.22     1.18    1.14    1.12     22
## 6 Antigu~   13300      40   South~    5.84     5.9     5.89    5.88     NA
## # ... with 12 more variables: cpi_2013 <dbl>, cpi_2014 <dbl>, cpi_2015 <dbl>,
## #   cpi_2016 <dbl>, cpi_2017 <dbl>, lifeexp_2012 <dbl>, lifeexp_2013 <dbl>,
## #   lifeexp_2014 <dbl>, lifeexp_2015 <dbl>, lifeexp_2016 <dbl>,
## #   lifeexp_2017 <dbl>, lifeexp_2018 <dbl>
```

Wide data have one row per unit while long data have more than one row per unit. To convert wide data to long, use `pivot_longer()`. There are several different ways to get the same output.

```
# Select columns using tidy-select
# Dictate pattern with names_sep
long_tib1 <- pivot_longer(tib1,
```

```

      contains("_"),
      names_to = c("var", "year"),
      names_sep = "_"

```

Select columns using column indices

```

pivot_longer(tib1,
             5:21,
             names_to = c("var", "year"),
             names_sep = "_")

```

A tibble: 3,315 x 7

```

##   country      gdp_2015 gini_2015 region      var  year  value
##   <chr>         <dbl>    <dbl> <chr>      <chr> <chr> <dbl>
## 1 Afghanistan    574      36.8 Asia & Pacific co2    2015  0.262
## 2 Afghanistan    574      36.8 Asia & Pacific co2    2016  0.245
## 3 Afghanistan    574      36.8 Asia & Pacific co2    2017  0.247
## 4 Afghanistan    574      36.8 Asia & Pacific co2    2018  0.254
## 5 Afghanistan    574      36.8 Asia & Pacific cpi     2012    8
## 6 Afghanistan    574      36.8 Asia & Pacific cpi     2013    8
## 7 Afghanistan    574      36.8 Asia & Pacific cpi     2014   12
## 8 Afghanistan    574      36.8 Asia & Pacific cpi     2015   11
## 9 Afghanistan    574      36.8 Asia & Pacific cpi     2016   15
## 10 Afghanistan   574      36.8 Asia & Pacific cpi     2017   15

```

... with 3,305 more rows

Dictate pattern with names_pattern

```

pivot_longer(tib1,
             5:21,
             names_to = c("var", "year"),
             names_pattern = "(.*)_(.*)")

```

A tibble: 3,315 x 7

```

##   country      gdp_2015 gini_2015 region      var  year  value
##   <chr>         <dbl>    <dbl> <chr>      <chr> <chr> <dbl>
## 1 Afghanistan    574      36.8 Asia & Pacific co2    2015  0.262
## 2 Afghanistan    574      36.8 Asia & Pacific co2    2016  0.245
## 3 Afghanistan    574      36.8 Asia & Pacific co2    2017  0.247
## 4 Afghanistan    574      36.8 Asia & Pacific co2    2018  0.254
## 5 Afghanistan    574      36.8 Asia & Pacific cpi     2012    8
## 6 Afghanistan    574      36.8 Asia & Pacific cpi     2013    8
## 7 Afghanistan    574      36.8 Asia & Pacific cpi     2014   12
## 8 Afghanistan    574      36.8 Asia & Pacific cpi     2015   11
## 9 Afghanistan    574      36.8 Asia & Pacific cpi     2016   15
## 10 Afghanistan   574      36.8 Asia & Pacific cpi     2017   15

```

... with 3,305 more rows

To go from long data to wide, use `pivot_wider`. There are several options to dictate the names of the newly created variables.

```

wide_tib1 <- pivot_wider(long_tib1,
                        names_from = c("var", "year"),
                        values_from = "value")
head(wide_tib1)

```

A tibble: 6 x 21

```

##   country region gdp_2015 gini_2015 co2_2015 co2_2016 co2_2017 co2_2018 cpi_2012

```



```
##   <chr>   <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 Afghan~ Asia ~      574       36.8      0.262      0.245      0.247      0.254       8
## 2 Albania Europe     4520       29       1.6       1.57      1.61      1.59      33
## 3 Algeria Arab ~     4780      27.6      3.8       3.64      3.56      3.69      34
## 4 Andorra Europe   42100      40      5.97      6.07      6.27      6.12      NA
## 5 Angola  Africa    3750     42.6     1.22      1.18      1.14      1.12      22
## 6 Antigu~ South~   13300     40     5.84      5.9       5.89      5.88      NA
## # ... with 12 more variables: cpi_2013 <dbl>, cpi_2014 <dbl>, cpi_2015 <dbl>,
## #   cpi_2016 <dbl>, cpi_2017 <dbl>, lifeexp_2012 <dbl>, lifeexp_2013 <dbl>,
## #   lifeexp_2014 <dbl>, lifeexp_2015 <dbl>, lifeexp_2016 <dbl>,
## #   lifeexp_2017 <dbl>, lifeexp_2018 <dbl>
```

```
pivot_wider(long_tib1,
             names_from = c("var", "year"),
             values_from = "value",
             names_sep = ".")
```

```
## # A tibble: 195 x 21
##   country      region      gdp.2015 gini.2015 co2.2015 co2.2016 co2.2017 co2.2018
##   <chr>        <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 Afghanistan Asia & Pa~      574       36.8      0.262      0.245      0.247      0.254
## 2 Albania     Europe     4520       29       1.6       1.57      1.61      1.59
## 3 Algeria     Arab Stat~  4780      27.6      3.8       3.64      3.56      3.69
## 4 Andorra     Europe   42100      40      5.97      6.07      6.27      6.12
## 5 Angola      Africa    3750     42.6     1.22      1.18      1.14      1.12
## 6 Antigua an~ South/Lat~ 13300     40     5.84      5.9       5.89      5.88
## 7 Argentina   South/Lat~ 10600     41.8     4.64      4.6       4.55      4.41
## 8 Armenia     Europe    3920     31.9     1.65      1.76      1.7       1.89
## 9 Australia   Asia & Pa~ 55100     32.3     16.8      17        17        16.9
## 10 Austria    Europe   47800     30.6     7.7       7.7       7.94      7.75
## # ... with 185 more rows, and 13 more variables: cpi.2012 <dbl>,
## #   cpi.2013 <dbl>, cpi.2014 <dbl>, cpi.2015 <dbl>, cpi.2016 <dbl>,
## #   cpi.2017 <dbl>, lifeexp.2012 <dbl>, lifeexp.2013 <dbl>, lifeexp.2014 <dbl>,
## #   lifeexp.2015 <dbl>, lifeexp.2016 <dbl>, lifeexp.2017 <dbl>,
## #   lifeexp.2018 <dbl>
```

It might be useful to reference [this chapter on strings and regular expressions](#). There are many ways to represent different patterns in character strings, and a standardized approach exists to minimize the need to type out everything explicitly.

Pipes

The `magrittr` package contains the pipe operator, `%>%`. The purpose of this operator is to make code clearer and more efficient. The idea is to minimize unnecessary saved objects. For example, if you are cleaning a dataset, it would be cumbersome to save a new data frame for each step in the cleaning process. Pipe operators, or pipes, help with this.

The idea is that the pipe forwards a value to the next function. The two lines result in the same output. The first argument of `filter()` is forwarded by the pipe operator.

```
filter(tib1, region == "North America")
```

```
## # A tibble: 2 x 21
##   country gdp_2015 gini_2015 region co2_2015 co2_2016 co2_2017 co2_2018 cpi_2012
##   <chr>      <dbl>      <dbl> <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 Canada    50300      31.7 North~     16       15.5      15.6      15.3       84
```

```
## 2 United~    52100      41.3 North~    16.9    16.4    16.2    16.6      73
## # ... with 12 more variables: cpi_2013 <dbl>, cpi_2014 <dbl>, cpi_2015 <dbl>,
## #   cpi_2016 <dbl>, cpi_2017 <dbl>, lifeexp_2012 <dbl>, lifeexp_2013 <dbl>,
## #   lifeexp_2014 <dbl>, lifeexp_2015 <dbl>, lifeexp_2016 <dbl>,
## #   lifeexp_2017 <dbl>, lifeexp_2018 <dbl>
```

```
tib1 %>% filter(region == "North America")
```

```
## # A tibble: 2 x 21
##   country gdp_2015 gini_2015 region co2_2015 co2_2016 co2_2017 co2_2018 cpi_2012
##   <chr>      <dbl>      <dbl> <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 Canada    50300      31.7 North~    16      15.5      15.6      15.3      84
## 2 United~    52100      41.3 North~    16.9    16.4      16.2      16.6      73
## # ... with 12 more variables: cpi_2013 <dbl>, cpi_2014 <dbl>, cpi_2015 <dbl>,
## #   cpi_2016 <dbl>, cpi_2017 <dbl>, lifeexp_2012 <dbl>, lifeexp_2013 <dbl>,
## #   lifeexp_2014 <dbl>, lifeexp_2015 <dbl>, lifeexp_2016 <dbl>,
## #   lifeexp_2017 <dbl>, lifeexp_2018 <dbl>
```

Pipe operators are especially useful when there are several operations being applied to the same object.

```
tib1 %>%
  distinct() %>%
  full_join(pop, by = "country") %>%
  arrange(desc(region), desc(population)) %>%
  head()
```

```
## # A tibble: 6 x 22
##   country gdp_2015 gini_2015 region co2_2015 co2_2016 co2_2017 co2_2018 cpi_2012
##   <chr>      <dbl>      <dbl> <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 Brazil    11400      51.6 South~    2.42      2.2      2.23      2.18      43
## 2 Mexico    10000      46.5 South~    3.96      3.94      3.95      3.79      34
## 3 Colomb~    7580      51.7 South~    1.96      2.01      1.91      1.96      36
## 4 Argent~   10600      41.8 South~    4.64      4.6      4.55      4.41      35
## 5 Peru      6110      43.7 South~    1.69      1.82      1.68      1.74      38
## 6 Venezu~    NA      46.9 South~    5.69      5.47      5.23      4.81      19
## # ... with 13 more variables: cpi_2013 <dbl>, cpi_2014 <dbl>, cpi_2015 <dbl>,
## #   cpi_2016 <dbl>, cpi_2017 <dbl>, lifeexp_2012 <dbl>, lifeexp_2013 <dbl>,
## #   lifeexp_2014 <dbl>, lifeexp_2015 <dbl>, lifeexp_2016 <dbl>,
## #   lifeexp_2017 <dbl>, lifeexp_2018 <dbl>, population <dbl>
```

To highlight the utility of pipe operators, consider these alternatives. They produce the same results. The first approach results in two objects that are not necessary for the final analysis, `tmp1` and `tmp2`. These objects are created with the sole purpose of being used in other functions. If the dataset is large, saving different versions of it can be burdensome. Additionally, the workspace becomes messy with so many temporary objects. While the second approach avoids temporary versions, it is difficult to read and understand.

```
tmp1 <- distinct(tib1)
tmp2 <- full_join(tmp1, pop, by = "country")
df <- arrange(tmp2, desc(region), desc(population))
head(df)
```

```
## # A tibble: 6 x 22
##   country gdp_2015 gini_2015 region co2_2015 co2_2016 co2_2017 co2_2018 cpi_2012
##   <chr>      <dbl>      <dbl> <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 Brazil    11400      51.6 South~    2.42      2.2      2.23      2.18      43
## 2 Mexico    10000      46.5 South~    3.96      3.94      3.95      3.79      34
## 3 Colomb~    7580      51.7 South~    1.96      2.01      1.91      1.96      36
```

```
## 4 Argent~    10600      41.8 South~    4.64    4.6    4.55    4.41    35
## 5 Peru      6110      43.7 South~    1.69    1.82    1.68    1.74    38
## 6 Venezu~    NA      46.9 South~    5.69    5.47    5.23    4.81    19
## # ... with 13 more variables: cpi_2013 <dbl>, cpi_2014 <dbl>, cpi_2015 <dbl>,
## #   cpi_2016 <dbl>, cpi_2017 <dbl>, lifeexp_2012 <dbl>, lifeexp_2013 <dbl>,
## #   lifeexp_2014 <dbl>, lifeexp_2015 <dbl>, lifeexp_2016 <dbl>,
## #   lifeexp_2017 <dbl>, lifeexp_2018 <dbl>, population <dbl>

head(arrange(distinct(full_join(tib1, pop, by = "country")), desc(region), desc(population)))

## # A tibble: 6 x 22
##   country gdp_2015 gini_2015 region co2_2015 co2_2016 co2_2017 co2_2018 cpi_2012
##   <chr>    <dbl>    <dbl> <chr>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 Brazil  11400      51.6 South~    2.42    2.2    2.23    2.18    43
## 2 Mexico  10000      46.5 South~    3.96    3.94    3.95    3.79    34
## 3 Colomb~   7580      51.7 South~    1.96    2.01    1.91    1.96    36
## 4 Argent~  10600      41.8 South~    4.64    4.6    4.55    4.41    35
## 5 Peru    6110      43.7 South~    1.69    1.82    1.68    1.74    38
## 6 Venezu~    NA      46.9 South~    5.69    5.47    5.23    4.81    19
## # ... with 13 more variables: cpi_2013 <dbl>, cpi_2014 <dbl>, cpi_2015 <dbl>,
## #   cpi_2016 <dbl>, cpi_2017 <dbl>, lifeexp_2012 <dbl>, lifeexp_2013 <dbl>,
## #   lifeexp_2014 <dbl>, lifeexp_2015 <dbl>, lifeexp_2016 <dbl>,
## #   lifeexp_2017 <dbl>, lifeexp_2018 <dbl>, population <dbl>
```


Note that the pipe operator can be used with functions outside of the tidyverse functions.

```
full_join(df, pop, by = "country") %>%
  write.csv("country_info.csv")
```

Pipe operators can forward objects to other arguments besides the first one. A period (.) indicates this. Here is an example with plotting (see chapter 5).

```
tib1 %>%
  distinct() %>%
  full_join(pop, by = "country") %>%
  qplot(x = gini_2015, data = .)
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



Getting comfortable with `%>%` can vastly smooth your workflow in R.

Practice Exercises 5.3

1. Use pipes to accomplish the following tasks on `tib1`: select `country`, `region`, `co2_2015`, and `co2_2016`, remove rows with missing values for either CO2 variables, create a variable that is `TRUE` when CO2 emissions in 2016 are smaller than those in 2015, and only keep the rows where this variable is `TRUE`.
2. Look at the documentation for `?magrittr`. There are four types of pipes. Take a moment to familiarize yourself with their differences.

Further Reading

There are many great resources online, including cheat sheets. [Here](#) is one for `dplyr`. Save this cheat sheet if you find it useful! More cheat sheets can be found [here](#).

The above information comes from chapters 5.1-5.3, 6, and 21 of Boehmke (2016), chapters 2.2.5 and 3 of Zamora Saiz et al. (2020). See Zamora Saiz et al. (2020) chapter 3 for information on `data.table`.

References

- Boehmke, Bradley C. 2016. *Data Wrangling with R*. Use R! Springer. <https://link-springer-com.proxy.lib.duke.edu/content/pdf/10.1007%2F978-3-319-45599-0.pdf>.
- Zamora Saiz, Alfonso, Carlos Quesada González, Lluís Hurtado Gil, and Diego Mondéjar Ruiz. 2020. *An Introduction to Data Analysis in R: Hands-on Coding, Data Mining, Visualization and Statistics from Scratch*.